

04 Performance Measures

2021-01-06

Contents

| | |
|---|-----------|
| Loading Libraries | 1 |
| Loading Data | 2 |
| Leader board Visualization | 4 |
| data preparation and Visulatization | 4 |
| Final tibble to visualize | 5 |
| Visualization | 5 |
| test performance with new data output from previous test | 9 |
| Grid Search | 10 |
| sort accoding to auc high to low | 11 |
| H2o Performance | 13 |
| load model | 13 |
| performance table | 19 |
| Visualize the trade of between the precision and the recall and the optimal threshold | 20 |
| Percision vs Recall plot | 22 |
| getting predictions_tbl from previous session | 23 |
| Gain and Lift calculations | 24 |
| Gain Plot | 25 |

Loading Libraries

```
library(tidyverse)
library(tidyquant)
library(broom)
library(umap)
```

```

library(readxl)
library(h2o)
library(readxl)
library(rsample)
library(recipes)
library(PerformanceAnalytics)
library(h2o)
library(cowplot)
library(glue)

library(h2o)
library(tidyquant)
library(broom)
library(umap)
library(tidyverse)
library(readxl)
library(rsample)
library(recipes)
library(PerformanceAnalytics) # for skewers

```

Loading Data

```

product_backorders_tbl <- read_csv("product_backorders.csv")

# Split into test and train
set.seed(seed = 1113)
split_obj <- rsample::initial_split(product_backorders_tbl, prop = 0.85)

# Assign training and test data
train_readable_tbl <- training(split_obj)
test_readable_tbl <- testing(split_obj)

recipe_obj <- recipe(went_on_backorder ~ ., data = train_readable_tbl) %>%
  step_zv(all_predictors()) %>%
  prep()

train_tbl <- bake(recipe_obj, new_data = train_readable_tbl)
test_tbl <- bake(recipe_obj, new_data = test_readable_tbl)

# set the predictor names
predictors <- c("national_inv", "lead_time", "forecast_3_month", "sales_3_month")

# specify the response
response <- "went_on_backorder"

h2o.init()

```

```
## Connection successful!
##
## R is connected to the H2O cluster:
##   H2O cluster uptime:      4 hours 42 minutes
##   H2O cluster timezone:    Europe/Berlin
##   H2O data parsing timezone: UTC
##   H2O cluster version:    3.32.0.1
##   H2O cluster version age: 2 months and 30 days
##   H2O cluster name:       H2O_started_from_R_ahmed_cgv301
##   H2O cluster total nodes: 1
##   H2O cluster total memory: 1.02 GB
##   H2O cluster total cores: 4
##   H2O cluster allowed cores: 4
##   H2O cluster healthy:    TRUE
##   H2O Connection ip:      localhost
##   H2O Connection port:    54321
##   H2O Connection proxy:   NA
##   H2O Internal Security:  FALSE
##   H2O API Extensions:     Amazon S3, Algos, AutoML, Core V3, TargetEncoder, Core V4
##   R Version:              R version 4.0.3 (2020-10-10)
```

```
# Split data into a training and a validation data frame
# Setting the seed is just for reproducibility
split_h2o <- h2o.splitFrame(as.h2o(train_tbl), ratios = c(0.85), seed = 1234)
```

```
## |
```

```
train_h2o <- split_h2o[[1]]
valid_h2o <- split_h2o[[2]]
test_h2o <- as.h2o(test_tbl)
```

```
## |
```

```
# Set the target and predictors
y <- response
x <- setdiff(names(train_h2o), y)

automl_models_h2o <- h2o.automl(
  x = x,
  y = y,
  training_frame = train_h2o,
  validation_frame = valid_h2o,
  leaderboard_frame = test_h2o,
  max_runtime_secs = 15,
  nfolds = 5
)
```

```
## |
```

```
## 03:17:11.229: User specified a validation frame with cross-validation still enabled. Please note that
## 03:17:11.230: AutoML: XGBoost is not available; skipping it. |
```

```
leaderboard_Model <- automl_models_h2o@leaderboard
```

```
Model<-automl_models_h2o@leader
```

```
h2o.saveModel(Model,path = "ml_journal-Automated_Machine_learning_2_model")
```

```
## [1] "D:\\Mechatronics_master\\Third semester\\Data science\\Machine learning\\ml_journal-AhmedShaheen"
```

Leader board Visualization

```
automl_models_h2o@leaderboard %>%
  as_tibble() %>%
  select(-c(mean_per_class_error, rmse, mse))
```

```
## # A tibble: 12 x 4
##   model_id                auc logloss aucpr
##   <chr>                  <dbl>   <dbl> <dbl>
## 1 StackedEnsemble_AllModels_AutoML_20210108_031711 0.938 0.200 0.696
## 2 StackedEnsemble_BestOfFamily_AutoML_20210108_031711 0.933 0.207 0.682
## 3 GBM_2_AutoML_20210108_031711 0.928 0.307 0.649
## 4 GBM_3_AutoML_20210108_031711 0.909 0.314 0.618
## 5 GBM_1_AutoML_20210108_031711 0.904 0.312 0.568
## 6 GBM_grid_1_AutoML_20210108_031711_model_1 0.887 0.286 0.529
## 7 GBM_4_AutoML_20210108_031711 0.869 0.324 0.585
## 8 DRF_1_AutoML_20210108_031711 0.854 0.692 0.525
## 9 XRT_1_AutoML_20210108_031711 0.760 3.01 0.387
## 10 GBM_5_AutoML_20210108_031711 0.741 0.347 0.325
## 11 DeepLearning_1_AutoML_20210108_031711 0.734 0.452 0.269
## 12 GLM_1_AutoML_20210108_031711 0.679 0.353 0.196
```

data preparation and Visulatization

```
data_transformed_tbl <- automl_models_h2o@leaderboard %>%
  as_tibble() %>%
  select(-c(aucpr, mean_per_class_error, rmse, mse)) %>%
  mutate(model_type = str_extract(model_id, "[^_]+")) %>%
  slice(1:15) %>%
  rownames_to_column(var = "rowname") %>%

  mutate(
    model_id = as_factor(model_id) %>% reorder(auc),
    model_type = as.factor(model_type)
  ) %>%

  pivot_longer(cols = -c(model_id, model_type, rowname),
```

```

      names_to = "key",
      values_to = "value",
      names_transform = list(key = forcats::fct_inorder)
    ) %>%
  mutate(model_id = paste0(rowname, ".", model_id) %>% as_factor() %>% fct_rev())

data_transformed_tbl

```

```

## # A tibble: 24 x 5
##   rowname model_id      model_type      key      value
##   <chr>    <fct>          <fct>      <fct>    <dbl>
## 1 1      1. StackedEnsemble_AllModels_AutoML_2021~ StackedEnsemb~ auc      0.938
## 2 1      1. StackedEnsemble_AllModels_AutoML_2021~ StackedEnsemb~ loglo~ 0.200
## 3 2      2. StackedEnsemble_BestOfFamily_AutoML_2~ StackedEnsemb~ auc      0.933
## 4 2      2. StackedEnsemble_BestOfFamily_AutoML_2~ StackedEnsemb~ loglo~ 0.207
## 5 3      3. GBM_2_AutoML_20210108_031711          GBM          auc      0.928
## 6 3      3. GBM_2_AutoML_20210108_031711          GBM          loglo~ 0.307
## 7 4      4. GBM_3_AutoML_20210108_031711          GBM          auc      0.909
## 8 4      4. GBM_3_AutoML_20210108_031711          GBM          loglo~ 0.314
## 9 5      5. GBM_1_AutoML_20210108_031711          GBM          auc      0.904
## 10 5     5. GBM_1_AutoML_20210108_031711          GBM          loglo~ 0.312
## # ... with 14 more rows

```

Final tibble to visualize

```

data_transformed_tbl

## # A tibble: 24 x 5
##   rowname model_id      model_type      key      value
##   <chr>    <fct>          <fct>      <fct>    <dbl>
## 1 1      1. StackedEnsemble_AllModels_AutoML_2021~ StackedEnsemb~ auc      0.938
## 2 1      1. StackedEnsemble_AllModels_AutoML_2021~ StackedEnsemb~ loglo~ 0.200
## 3 2      2. StackedEnsemble_BestOfFamily_AutoML_2~ StackedEnsemb~ auc      0.933
## 4 2      2. StackedEnsemble_BestOfFamily_AutoML_2~ StackedEnsemb~ loglo~ 0.207
## 5 3      3. GBM_2_AutoML_20210108_031711          GBM          auc      0.928
## 6 3      3. GBM_2_AutoML_20210108_031711          GBM          loglo~ 0.307
## 7 4      4. GBM_3_AutoML_20210108_031711          GBM          auc      0.909
## 8 4      4. GBM_3_AutoML_20210108_031711          GBM          loglo~ 0.314
## 9 5      5. GBM_1_AutoML_20210108_031711          GBM          auc      0.904
## 10 5     5. GBM_1_AutoML_20210108_031711          GBM          loglo~ 0.312
## # ... with 14 more rows

```

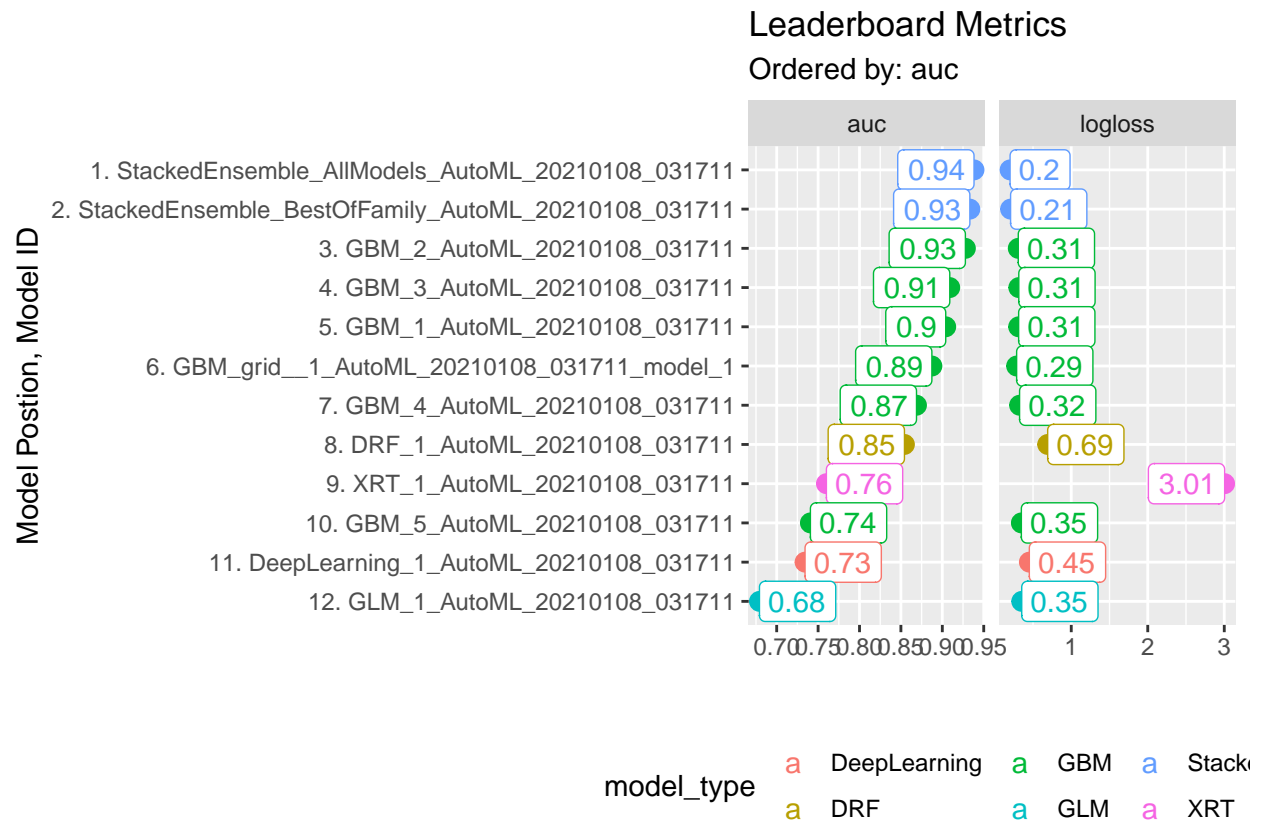
Visualization

```

data_transformed_tbl %>%
  ggplot(aes(value, model_id, color = model_type)) +
  geom_point(size = 3) +
  geom_label(aes(label = round(value, 2), hjust = "inward")) +

```

```
# Facet to break out logloss and auc
facet_wrap(~ key, scales = "free_x") +
labs(title = "Leaderboard Metrics",
      subtitle = paste0("Ordered by: ", "auc"),
      y = "Model Position, Model ID", x = "") +
theme(legend.position = "bottom")
```



```
# Grid Search
```

```
grid_search_model <- h2o.loadModel("ml_journal-Automated_Machine_learning_2_model/StackedEnsemble_AllModels_AutoML_20210107_224531")
```

```
grid_search_model
```

```
## Model Details:
```

```
## =====
```

```
##
```

```
## H2OBinomialModel: stackedensemble
```

```
## Model ID: StackedEnsemble_AllModels_AutoML_20210107_224531
```

```
## Number of Base Models: 10
```

```
##
```

```
## Base Models (count by algorithm type):
```

```
##
```

```
## deeplearning                drf                gbm                glm
```

```
##                1                2                6                1
```

```
##
```

```

## Metalearner:
##
## Metalearner algorithm: glm
## Metalearner cross-validation fold assignment:
##   Fold assignment scheme: AUTO
##   Number of folds: 5
##   Fold column: NULL
## Metalearner hyperparameters:
##
##
## H2OBinomialMetrics: stackedensemble
## ** Reported on training data. **
##
## MSE: 0.04664896
## RMSE: 0.2159837
## LogLoss: 0.1665542
## Mean Per-Class Error: 0.1411272
## AUC: 0.9695554
## AUCPR: 0.8361613
## Gini: 0.9391108
##
## Confusion Matrix (vertical: actual; across: predicted) for F1-optimal threshold:
##      No  Yes  Error  Rate
## No    8521  284 0.032254  =284/8805
## Yes    303  909 0.250000  =303/1212
## Totals 8824 1193 0.058600  =587/10017
##
## Maximum Metrics: Maximum metrics at their respective thresholds
##
##      metric threshold  value idx
## 1      max f1  0.256367  0.755925 212
## 2      max f2  0.135270  0.821907 266
## 3      max f0point5 0.424893  0.789763 157
## 4      max accuracy 0.293859  0.941400 199
## 5      max precision 0.997190  1.000000 0
## 6      max recall 0.027349  1.000000 378
## 7      max specificity 0.997190  1.000000 0
## 8      max absolute_mcc 0.175042  0.723872 248
## 9      max min_per_class_accuracy 0.124218  0.911414 272
## 10     max mean_per_class_accuracy 0.114705  0.913524 277
## 11     max tns 0.997190 8805.000000 0
## 12     max fns 0.997190 1201.000000 0
## 13     max fps 0.019690 8805.000000 399
## 14     max tps 0.027349 1212.000000 378
## 15     max tnr 0.997190 1.000000 0
## 16     max fnr 0.997190 0.990924 0
## 17     max fpr 0.019690 1.000000 399
## 18     max tpr 0.027349 1.000000 378
##
## Gains/Lift Table: Extract with 'h2o.gainsLift(<model>, <data>)' or 'h2o.gainsLift(<model>, valid=<T/I>)'
## H2OBinomialMetrics: stackedensemble
## ** Reported on validation data. **
##
## MSE: 0.06603241
## RMSE: 0.2569677

```

```

## LogLoss: 0.2243434
## Mean Per-Class Error: 0.1681104
## AUC: 0.925822
## AUCPR: 0.6400195
## Gini: 0.8516441
##
## Confusion Matrix (vertical: actual; across: predicted) for F1-optimal threshold:
##      No Yes  Error  Rate
## No    1930 174 0.082700 =174/2104
## Yes     72 212 0.253521 =72/284
## Totals 2002 386 0.103015 =246/2388
##
## Maximum Metrics: Maximum metrics at their respective thresholds
##      metric threshold  value idx
## 1      max f1 0.179216 0.632836 221
## 2      max f2 0.112725 0.744298 260
## 3      max f0point5 0.482100 0.625000 111
## 4      max accuracy 0.482100 0.910804 111
## 5      max precision 0.997476 1.000000 0
## 6      max recall 0.024669 1.000000 385
## 7      max specificity 0.997476 1.000000 0
## 8      max absolute_mcc 0.179216 0.583687 221
## 9      max min_per_class_accuracy 0.114479 0.866197 259
## 10     max mean_per_class_accuracy 0.112725 0.869605 260
## 11     max tns 0.997476 2104.000000 0
## 12     max fns 0.997476 282.000000 0
## 13     max fps 0.019476 2104.000000 399
## 14     max tps 0.024669 284.000000 385
## 15     max tnr 0.997476 1.000000 0
## 16     max fnr 0.997476 0.992958 0
## 17     max fpr 0.019476 1.000000 399
## 18     max tpr 0.024669 1.000000 385
##
## Gains/Lift Table: Extract with 'h2o.gainsLift(<model>, <data>)' or 'h2o.gainsLift(<model>, valid=<T/
## H2OBinoMialMetrics: stackedensemble
## ** Reported on cross-validation data. **
## ** 5-fold cross-validation on training data (Metrics computed for combined holdout predictions) **
##
## MSE: 0.06744484
## RMSE: 0.2597014
## LogLoss: 0.2337148
## Mean Per-Class Error: 0.2156106
## AUC: 0.9138266
## AUCPR: 0.623122
## Gini: 0.8276533
##
## Confusion Matrix (vertical: actual; across: predicted) for F1-optimal threshold:
##      No Yes  Error  Rate
## No    11528 632 0.051974 =632/12160
## Yes     625 1023 0.379248 =625/1648
## Totals 12153 1655 0.091034 =1257/13808
##
## Maximum Metrics: Maximum metrics at their respective thresholds
##      metric threshold  value idx

```



```
## 1          max f1 0.255882    0.619437 212
## 2          max f2 0.078528    0.698761 306
## 3          max f0point5 0.331120    0.632699 186
## 4          max accuracy 0.331120    0.912080 186
## 5          max precision 0.999404    0.937500 0
## 6          max recall 0.017271    1.000000 387
## 7          max specificity 0.999404    0.999836 0
## 8          max absolute_mcc 0.255882    0.567738 212
## 9  max min_per_class_accuracy 0.086554    0.838816 300
## 10 max mean_per_class_accuracy 0.077422    0.842540 307
## 11          max tns 0.999404 12158.000000 0
## 12          max fns 0.999404 1618.000000 0
## 13          max fps 0.008532 12160.000000 399
## 14          max tps 0.017271 1648.000000 387
## 15          max tnr 0.999404    0.999836 0
## 16          max fnr 0.999404    0.981796 0
## 17          max fpr 0.008532    1.000000 399
## 18          max tpr 0.017271    1.000000 387
##
## Gains/Lift Table: Extract with 'h2o.gainsLift(<model>, <data>)' or 'h2o.gainsLift(<model>, valid=<T/
```

```
test_tbl
```

```
## # A tibble: 2,857 x 23
##   sku national_inv lead_time in_transit_qty forecast_3_month
##   <dbl>         <dbl>    <dbl>         <dbl>         <dbl>
## 1 1.12e6           7         8             0             0
## 2 1.12e6           0        12             0            163
## 3 1.13e6        150         8             0            325
## 4 1.13e6           0         8             0             3
## 5 1.13e6           3         3             0            16
## 6 1.13e6        -62         8             0            120
## 7 1.13e6           0         8             0             2
## 8 1.13e6          11         2            10            27
## 9 1.13e6           1         2             4             4
## 10 1.14e6          -1         8             0             9
## # ... with 2,847 more rows, and 18 more variables: forecast_6_month <dbl>,
## #   forecast_9_month <dbl>, sales_1_month <dbl>, sales_3_month <dbl>,
## #   sales_6_month <dbl>, sales_9_month <dbl>, min_bank <dbl>,
## #   potential_issue <fct>, pieces_past_due <dbl>, perf_6_month_avg <dbl>,
## #   perf_12_month_avg <dbl>, local_bo_qty <dbl>, deck_risk <fct>,
## #   oe_constraint <fct>, ppap_risk <fct>, stop_auto_buy <fct>, rev_stop <fct>,
## #   went_on_backorder <fct>
```

test performance with new data output from previous test

```
h2o.performance(grid_search_model, newdata = as.h2o(test_tbl))
```

```
## |
```

```
|
```

```

## H2OBinomialMetrics: stackedensemble
##
## MSE: 0.06315357
## RMSE: 0.2513037
## LogLoss: 0.2183313
## Mean Per-Class Error: 0.1711951
## AUC: 0.9297926
## AUCPR: 0.6499886
## Gini: 0.8595852
##
## Confusion Matrix (vertical: actual; across: predicted) for F1-optimal threshold:
##      No Yes  Error  Rate
## No    2339 184 0.072929 =184/2523
## Yes     90 244 0.269461  =90/334
## Totals 2429 428 0.095905 =274/2857
##
## Maximum Metrics: Maximum metrics at their respective thresholds
##      metric threshold  value idx
## 1      max f1 0.201687 0.640420 221
## 2      max f2 0.096409 0.728395 284
## 3      max f0point5 0.389786 0.656385 154
## 4      max accuracy 0.389786 0.918446 154
## 5      max precision 0.996593 1.000000 0
## 6      max recall 0.031624 1.000000 370
## 7      max specificity 0.996593 1.000000 0
## 8      max absolute_mcc 0.187938 0.593279 227
## 9      max min_per_class_accuracy 0.106826 0.853293 275
## 10     max mean_per_class_accuracy 0.096409 0.863535 284
## 11     max tns 0.996593 2523.000000 0
## 12     max fns 0.996593 332.000000 0
## 13     max fps 0.020097 2523.000000 399
## 14     max tps 0.031624 334.000000 370
## 15     max tnr 0.996593 1.000000 0
## 16     max fnr 0.996593 0.994012 0
## 17     max fpr 0.020097 1.000000 399
## 18     max tpr 0.031624 1.000000 370
##
## Gains/Lift Table: Extract with 'h2o.gainsLift(<model>, <data>)' or 'h2o.gainsLift(<model>, valid=<T/I/

```

Grid Search

```

search_grid_01 <- h2o.grid(
  algorithm = "deeplearning",
  grid_id = "search_grid_01",
  x = x,
  y = y,

  # training and validation frame and crossfold validation
  training_frame = train_h2o,
  validation_frame = valid_h2o,
  nfolds = 5,

```

```

# Hyperparameters: Use deeplearning_h2o@allparameters to see all
hyper_params = list(
    # Use some combinations (the first one was the original)
    hidden = list(c(10, 10, 10), c(50, 20, 10), c(20, 20, 20)),
    epochs = c(10, 50, 20)
)

```

```
## |
```

```
search_grid_01
```

```

## H2O Grid Details
## =====
##
## Grid ID: search_grid_01
## Used hyper parameters:
##   - epochs
##   - hidden
## Number of models: 108
## Number of failed models: 0
##
## Hyper-Parameter Search Summary: ordered by increasing logloss
##           epochs      hidden      model_ids      logloss
## 1 10.401183024268088 [50, 20, 10] search_grid_01_model_94 0.593392794267379
## 2 10.396790876881736 [10, 10, 10] search_grid_01_model_19 0.6872272269441417
## 3 10.39353646494803 [10, 10, 10] search_grid_01_model_91 0.7753158814829997
## 4 10.390115491579332 [10, 10, 10] search_grid_01_model_10 0.8019901718186699
## 5 10.379148529101434 [50, 20, 10] search_grid_01_model_76 0.8212459493715552
##
## ---
##           epochs      hidden      model_ids      logloss
## 103 20.79287977545045 [20, 20, 20] search_grid_01_model_72 3.982996617678043
## 104 52.01446341563045 [20, 20, 20] search_grid_01_model_62 3.999608770441817
## 105 52.00488790148506 [10, 10, 10] search_grid_01_model_56 4.741461845786739
## 106 52.012464918336036 [20, 20, 20] search_grid_01_model_98 4.87022132619315
## 107 20.79885310027837 [50, 20, 10] search_grid_01_model_15 5.1541284744691325
## 108 51.99653962675832 [20, 20, 20] search_grid_01_model_71 6.757877764291338

```

sort accoding to auc high to low

```
h2o.getGrid(grid_id = "search_grid_01", sort_by = "auc", decreasing = TRUE)
```

```

## H2O Grid Details
## =====
##
## Grid ID: search_grid_01
## Used hyper parameters:
##   - epochs

```

```
## - hidden
## Number of models: 108
## Number of failed models: 0
##
## Hyper-Parameter Search Summary: ordered by decreasing auc
##           epochs      hidden      model_ids      auc
## 1  51.98202355143529 [50, 20, 10] search_grid_01_model_50 0.6766130247588783
## 2  10.401183024268088 [50, 20, 10] search_grid_01_model_94 0.6759369161583418
## 3   20.78583674094471 [50, 20, 10] search_grid_01_model_78  0.65798493289314
## 4  20.798433494599845 [50, 20, 10] search_grid_01_model_24 0.6345598332907512
## 5   10.40730809032598 [20, 20, 20] search_grid_01_model_16  0.628448158852836
##
## ---
##           epochs      hidden      model_ids      auc
## 103 10.39973847867544 [50, 20, 10] search_grid_01_model_4 0.5210769333642055
## 104 51.993617845941706 [10, 10, 10] search_grid_01_model_38 0.5207654263940342
## 105 52.002001061293505 [10, 10, 10] search_grid_01_model_92 0.5156528198055059
## 106 20.785329625967172 [10, 10, 10] search_grid_01_model_57 0.5130598392788708
## 107 20.79758024534127 [20, 20, 20] search_grid_01_model_45 0.5100362381036024
## 108 20.802074059358134 [20, 20, 20] search_grid_01_model_90 0.5007095672186382
```

```
search_grid_01_model_1 <- h2o.getModel("search_grid_01_model_1")
search_grid_01_model_1 %>% h2o.auc(train = T, valid = T, xval = T)
```

```
##      train      valid      xval
## 0.6521814 0.6270526 0.5897518
```

```
search_grid_01_model_1 %>%
  h2o.performance(newdata = as.h2o(test_tbl))
```

```
## |

## H2OBinomialMetrics: deeplearning
##
## MSE: 0.1052787
## RMSE: 0.3244668
## LogLoss: 0.3974823
## Mean Per-Class Error: 0.404194
## AUC: 0.6518983
## AUCPR: 0.2197715
## Gini: 0.3037967
##
## Confusion Matrix (vertical: actual; across: predicted) for F1-optimal threshold:
##           No Yes      Error      Rate
## No      1964 559 0.221562 =559/2523
## Yes      196 138 0.586826 =196/334
## Totals 2160 697 0.264263 =755/2857
##
## Maximum Metrics: Maximum metrics at their respective thresholds
##           metric threshold      value idx
## 1           max f1 0.158513 0.267701 206
## 2           max f2 0.040741 0.445923 350
```

```
## 3          max f0point5 0.378959    0.266749  74
## 4          max accuracy 0.932781    0.885194  10
## 5          max precision 0.932781    0.607143  10
## 6          max recall   0.000034    1.000000 399
## 7          max specificity 0.999862    0.998018  0
## 8          max absolute_mcc 0.783253    0.166765  22
## 9  max min_per_class_accuracy 0.102825    0.577844 269
## 10 max mean_per_class_accuracy 0.050803    0.603219 333
## 11          max tns 0.999862 2518.000000  0
## 12          max fns 0.999862  328.000000  0
## 13          max fps 0.000034 2523.000000 399
## 14          max tps 0.000034  334.000000 399
## 15          max tnr 0.999862    0.998018  0
## 16          max fnr 0.999862    0.982036  0
## 17          max fpr 0.000034    1.000000 399
## 18          max tpr 0.000034    1.000000 399
##
## Gains/Lift Table: Extract with 'h2o.gainsLift(<model>, <data>)' or 'h2o.gainsLift(<model>, valid=<T/
```

H2o Performance

load model

```
stacked_ensemble_h2o <- h2o.loadModel("ml_journal-Automated_Machine_learning_2_model/StackedEnsemble_Algorithm")
performance_h2o <- h2o.performance(stacked_ensemble_h2o, newdata = as.h2o(test_tbl))
```

```
##      |
```

```
typeof(performance_h2o)
```

```
## [1] "S4"
```

```
performance_h2o %>% slotNames()
```

```
## [1] "algorithm" "on_train" "on_valid" "on_xval" "metrics"
```

```
performance_h2o@metrics
```

```
## $model
## $model$'__meta'
## $model$'__meta'$schema_version
## [1] 3
##
## $model$'__meta'$schema_name
## [1] "ModelKeyV3"
##
## $model$'__meta'$schema_type
```

```

## [1] "Key<Model>"
##
##
## $model$name
## [1] "StackedEnsemble_AllModels_AutoML_20210108_015304"
##
## $model$type
## [1] "Key<Model>"
##
## $model$URL
## [1] "/3/Models/StackedEnsemble_AllModels_AutoML_20210108_015304"
##
##
## $model_checksum
## [1] "-2137875469075093584"
##
## $frame
## $frame$name
## [1] "test_tbl_sid_90af_462"
##
##
## $frame_checksum
## [1] "-1590291148178685696"
##
## $description
## NULL
##
## $scoring_time
## [1] 1.610073e+12
##
## $predictions
## NULL
##
## $MSE
## [1] 0.05601617
##
## $RMSE
## [1] 0.2366774
##
## $nobs
## [1] 2857
##
## $custom_metric_name
## NULL
##
## $custom_metric_value
## [1] 0
##
## $r2
## [1] 0.457412
##
## $logloss
## [1] 0.196452
##

```

```

## $AUC
## [1] 0.9417586
##
## $pr_auc
## [1] 0.7119371
##
## $Gini
## [1] 0.8835172
##
## $mean_per_class_error
## [1] 0.1758813
##
## $domain
## [1] "No" "Yes"
##
## $cm
## $cm$`__meta`
## $cm$`__meta`$schema_version
## [1] 3
##
## $cm$`__meta`$schema_name
## [1] "ConfusionMatrixV3"
##
## $cm$`__meta`$schema_type
## [1] "ConfusionMatrix"
##
##
## $cm$table
## Confusion Matrix: Row labels: Actual class; Column labels: Predicted class
##           No Yes  Error      Rate
## No      2406 117 0.0464 = 117 / 2,523
## Yes      102 232 0.3054 = 102 / 334
## Totals 2508 349 0.0767 = 219 / 2,857
##
##
## $thresholds_and_metric_scores
## Metrics for Thresholds: Binomial metrics as a function of classification thresholds
##   threshold      f1      f2 f0point5 accuracy precision  recall specificity
## 1  0.998898 0.011905 0.007474 0.029240 0.883794 1.000000 0.005988 1.000000
## 2  0.996792 0.029499 0.018643 0.070621 0.884844 1.000000 0.014970 1.000000
## 3  0.995788 0.058140 0.037147 0.133690 0.886594 1.000000 0.029940 1.000000
## 4  0.994837 0.074928 0.048184 0.168394 0.887644 1.000000 0.038922 1.000000
## 5  0.993040 0.096591 0.062777 0.209360 0.888694 0.944444 0.050898 0.999604
##   absolute_mcc min_per_class_accuracy mean_per_class_accuracy  tns fns fps tps
## 1    0.072744                0.005988          0.502994 2523 332  0  2
## 2    0.115079                0.014970          0.507485 2523 329  0  5
## 3    0.162889                0.029940          0.514970 2523 324  0 10
## 4    0.185820                0.038922          0.519461 2523 321  0 13
## 5    0.205079                0.050898          0.525251 2522 317  1 17
##           tnr      fnr      fpr      tpr idx
## 1 1.000000 0.994012 0.000000 0.005988  0
## 2 1.000000 0.985030 0.000000 0.014970  1
## 3 1.000000 0.970060 0.000000 0.029940  2
## 4 1.000000 0.961078 0.000000 0.038922  3

```

```

## 5 0.999604 0.949102 0.000396 0.050898 4
##
## ---
##      threshold      f1      f2 f0point5 accuracy precision  recall
## 395 0.020644 0.214860 0.406227 0.146056 0.145607 0.120360 1.000000
## 396 0.020348 0.213282 0.403967 0.144890 0.137557 0.119371 1.000000
## 397 0.020043 0.211526 0.401442 0.143594 0.128456 0.118272 1.000000
## 398 0.019719 0.210129 0.399426 0.142564 0.121106 0.117399 1.000000
## 399 0.019222 0.209667 0.398758 0.142224 0.118656 0.117111 1.000000
## 400 0.018879 0.209339 0.398283 0.141983 0.116906 0.116906 1.000000
##      specificity absolute_mcc min_per_class_accuracy mean_per_class_accuracy tns
## 395 0.032501 0.062545 0.032501 0.516250 82
## 396 0.023385 0.052834 0.023385 0.511692 59
## 397 0.013080 0.039331 0.013080 0.506540 33
## 398 0.004756 0.023630 0.004756 0.502378 12
## 399 0.001982 0.015234 0.001982 0.500991 5
## 400 0.000000 0.000000 0.000000 0.500000 0
##      fns fps tps      tnr      fnr      fpr      tpr idx
## 395 0 2441 334 0.032501 0.000000 0.967499 1.000000 394
## 396 0 2464 334 0.023385 0.000000 0.976615 1.000000 395
## 397 0 2490 334 0.013080 0.000000 0.986920 1.000000 396
## 398 0 2511 334 0.004756 0.000000 0.995244 1.000000 397
## 399 0 2518 334 0.001982 0.000000 0.998018 1.000000 398
## 400 0 2523 334 0.000000 0.000000 1.000000 1.000000 399
##
## $max_criteria_and_metric_scores
## Maximum Metrics: Maximum metrics at their respective thresholds
##      metric threshold      value idx
## 1      max f1 0.323580 0.679356 188
## 2      max f2 0.125719 0.745243 270
## 3      max f0point5 0.520073 0.705795 135
## 4      max accuracy 0.480941 0.927896 144
## 5      max precision 0.998898 1.000000 0
## 6      max recall 0.025434 1.000000 378
## 7      max specificity 0.998898 1.000000 0
## 8      max absolute_mcc 0.323580 0.636047 188
## 9      max min_per_class_accuracy 0.087553 0.865636 294
## 10     max mean_per_class_accuracy 0.071412 0.872407 306
## 11     max tns 0.998898 2523.000000 0
## 12     max fns 0.998898 332.000000 0
## 13     max fps 0.018879 2523.000000 399
## 14     max tps 0.025434 334.000000 378
## 15     max tnr 0.998898 1.000000 0
## 16     max fnr 0.998898 0.994012 0
## 17     max fpr 0.018879 1.000000 399
## 18     max tpr 0.025434 1.000000 378
##
## $gains_lift_table
## Gains/Lift Table: Avg response rate: 11.69 %, avg score: 12.70 %
##      group cumulative_data_fraction lower_threshold lift cumulative_lift
## 1      1      0.01015051 0.983449 8.258930 8.258930
## 2      2      0.02030102 0.949574 6.784121 7.521526
## 3      3      0.03010151 0.896397 6.109923 7.061934
## 4      4      0.04025201 0.854943 7.963969 7.289404

```



```

## 5      5      0.05005250      0.803196 5.498931      6.938822
## 6      6      0.10010501      0.429004 5.024664      5.981743
## 7      7      0.15015751      0.211347 3.170324      5.044603
## 8      8      0.20021001      0.115391 1.794523      4.232083
## 9      9      0.29996500      0.046491 0.930423      3.134098
## 10     10     0.40007000      0.034363 0.358905      2.439693
## 11     11     0.50017501      0.029974 0.029909      1.957399
## 12     12     0.59993000      0.027532 0.090041      1.646899
## 13     13     0.70003500      0.025604 0.089726      1.424223
## 14     14     0.79978999      0.023861 0.030014      1.250328
## 15     15     0.89989499      0.022050 0.000000      1.111241
## 16     16     1.00000000      0.018589 0.000000      1.000000
##      response_rate      score cumulative_response_rate cumulative_score
## 1      0.965517 0.992662      0.965517      0.992662
## 2      0.793103 0.971039      0.879310      0.981850
## 3      0.714286 0.921029      0.825581      0.962048
## 4      0.931034 0.880146      0.852174      0.941395
## 5      0.642857 0.829368      0.811189      0.919459
## 6      0.587413 0.621150      0.699301      0.770305
## 7      0.370629 0.310652      0.589744      0.617087
## 8      0.209790 0.158233      0.494755      0.502374
## 9      0.108772 0.069594      0.366394      0.358450
## 10     0.041958 0.039293      0.285214      0.278591
## 11     0.003497 0.031784      0.228831      0.229195
## 12     0.010526 0.028698      0.192532      0.195857
## 13     0.010490 0.026513      0.166500      0.171641
## 14     0.003509 0.024746      0.146171      0.153319
## 15     0.000000 0.022992      0.129911      0.138821
## 16     0.000000 0.020904      0.116906      0.127017
##      capture_rate cumulative_capture_rate      gain cumulative_gain
## 1      0.083832      0.083832 725.893042      725.893042
## 2      0.068862      0.152695 578.412141      652.152591
## 3      0.059880      0.212575 510.992301      606.193427
## 4      0.080838      0.293413 696.396861      628.940380
## 5      0.053892      0.347305 449.893071      593.882166
## 6      0.251497      0.598802 402.466396      498.174281
## 7      0.158683      0.757485 217.032369      404.460310
## 8      0.089820      0.847305 79.452284      323.208304
## 9      0.092814      0.940120 -6.957664      213.409820
## 10     0.035928      0.976048 -64.109543      143.969279
## 11     0.002994      0.979042 -97.009129      95.739871
## 12     0.008982      0.988024 -90.995903      64.689873
## 13     0.008982      0.997006 -91.027386      42.422305
## 14     0.002994      1.000000 -96.998634      25.032823
## 15     0.000000      1.000000 -100.000000      11.124076
## 16     0.000000      1.000000 -100.000000      0.000000
##      kolmogorov_smirnov
## 1      0.083436
## 2      0.149920
## 3      0.206630
## 4      0.286675
## 5      0.336604
## 6      0.564716
## 7      0.687727

```

```
## 8          0.732759
## 9          0.724900
## 10         0.652227
## 11         0.542260
## 12         0.439471
## 13         0.336285
## 14         0.226714
## 15         0.113357
## 16         0.000000
##
## $residual_deviance
## [1] 1122.527
##
## $null_deviance
## [1] 2061.286
##
## $AIC
## [1] 1138.527
##
## $null_degrees_of_freedom
## [1] 2856
##
## $residual_degrees_of_freedom
## [1] 2849
```

```
# Classifier Summary Metrics
```

```
h2o.auc(performance_h2o, train = T, valid = T, xval = T)
```

```
## [1] 0.9417586
```

```
# our value is [1] 0.9037603
```

```
h2o.auc(stacked_ensemble_h2o, train = T, valid = T, xval = T)
```

```
##      train      valid      xval
## 0.9787684 0.9548814 0.9304988
```

```
#      train      valid      xval
# 0.9320589 0.8932458 0.8576325
```

```
h2o.giniCoef(performance_h2o)
```

```
## [1] 0.8835172
```

```
# [1] 0.8075205
```

```
h2o.logloss(performance_h2o)
```

```
## [1] 0.196452
```

```
# [1] 0.2362433
```

```
# result for the training data
```

```
h2o.confusionMatrix(stacked_ensemble_h2o)
```

```
## Confusion Matrix (vertical: actual; across: predicted) for max f1 @ threshold = 0.341126899054185:
```

```
##      No Yes  Error   Rate
## No   8579 237 0.026883 =237/8816
## Yes   227 974 0.189009 =227/1201
## Totals 8806 1211 0.046321 =464/10017
```

```
# Confusion Matrix (vertical: actual; across: predicted) for max f1 @ threshold = 0.268791256959701
```

```
# result for the hold out set
```

```
h2o.confusionMatrix(performance_h2o)
```

```
## Confusion Matrix (vertical: actual; across: predicted) for max f1 @ threshold = 0.323579611938672:
```

```
##      No Yes  Error   Rate
## No   2406 117 0.046373 =117/2523
## Yes   102 232 0.305389 =102/334
## Totals 2508 349 0.076654 =219/2857
```

```
# Confusion Matrix (vertical: actual; across: predicted) for max f1 @ threshold = 0.25817823725732
```

performance table

```
performance_tbl <- performance_h2o %>%
  h2o.metric() %>%
  as.tibble()
```

```
performance_tbl %>%
  glimpse()
```

```
## Rows: 400
## Columns: 20
## $ threshold      <dbl> 0.9988981, 0.9967924, 0.9957879, 0.9948373,...
## $ f1              <dbl> 0.01190476, 0.02949853, 0.05813953, 0.07492...
## $ f2              <dbl> 0.007473842, 0.018642804, 0.037147103, 0.04...
## $ f0point5        <dbl> 0.02923977, 0.07062147, 0.13368984, 0.16839...
## $ accuracy        <dbl> 0.8837942, 0.8848442, 0.8865943, 0.8876444,...
## $ precision       <dbl> 1.0000000, 1.0000000, 1.0000000, 1.0000000,...
## $ recall          <dbl> 0.005988024, 0.014970060, 0.029940120, 0.03...
## $ specificity      <dbl> 1.0000000, 1.0000000, 1.0000000, 1.0000000,...
## $ absolute_mcc     <dbl> 0.07274403, 0.11507888, 0.16288896, 0.18581...
## $ min_per_class_accuracy <dbl> 0.005988024, 0.014970060, 0.029940120, 0.03...
## $ mean_per_class_accuracy <dbl> 0.5029940, 0.5074850, 0.5149701, 0.5194611,...
## $ tns             <dbl> 2523, 2523, 2523, 2523, 2522, 2522, 2522, 2...
## $ fns             <dbl> 332, 329, 324, 321, 317, 313, 312, 311, 309...
## $ fps             <dbl> 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 2, 3, 4, 4, 4...
```

```
## $ tps          <dbl> 2, 5, 10, 13, 17, 21, 22, 23, 25, 27, 29, 3...
## $ tnr          <dbl> 1.0000000, 1.0000000, 1.0000000, 1.0000000,...
## $ fnr          <dbl> 0.9940120, 0.9850299, 0.9700599, 0.9610778,...
## $ fpr          <dbl> 0.0000000000, 0.0000000000, 0.0000000000, 0...
## $ tpr          <dbl> 0.005988024, 0.014970060, 0.029940120, 0.03...
## $ idx          <int> 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 1...
```

```
# save our theme
theme_new <- theme(
  legend.position = "bottom",
  legend.key      = element_blank(),
  panel.background = element_rect(fill = "transparent"),
  panel.border     = element_rect(color = "black", fill = NA, size = 0.5),
  panel.grid.major = element_line(color = "grey", size = 0.333)
)
```

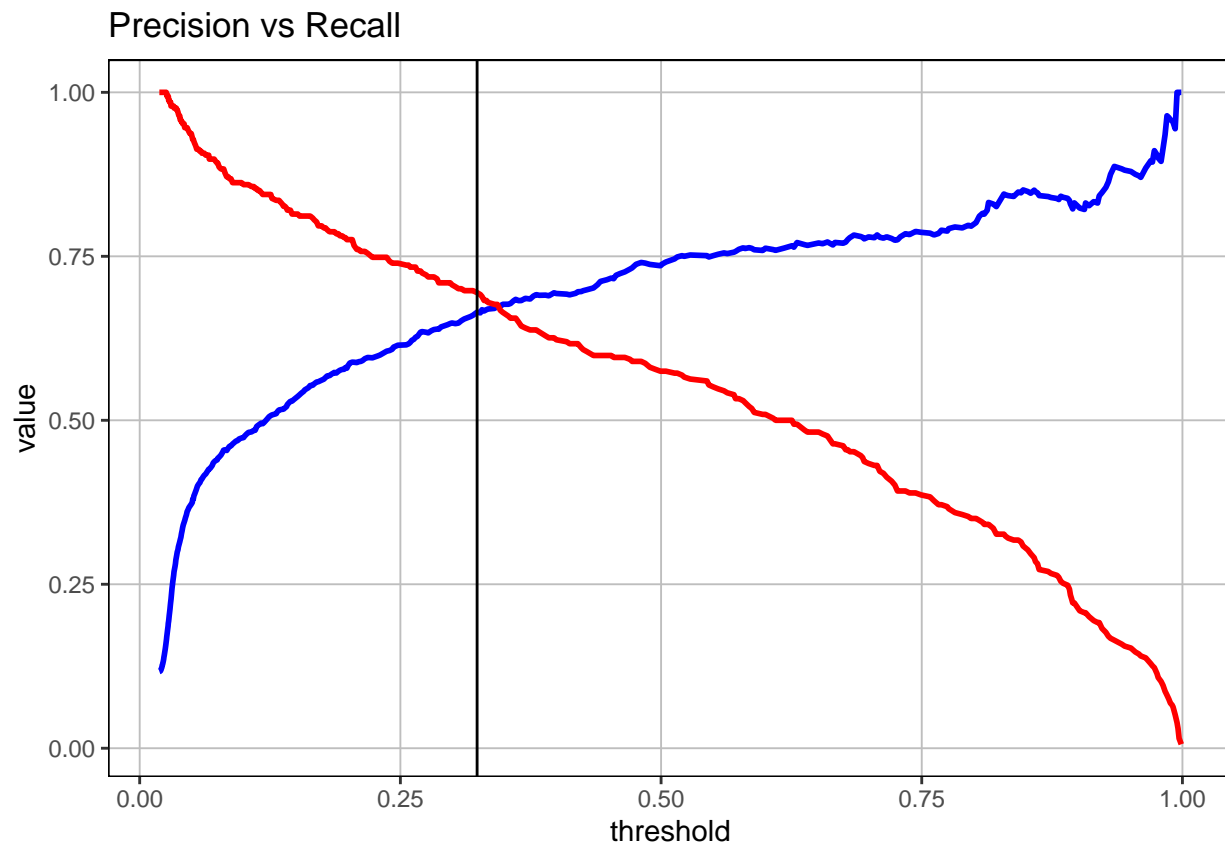
Visualize the trade of between the precision and the recall and the optimal threshold

```
performance_tbl %>%
  filter(f1 == max(f1))
```

```
## # A tibble: 1 x 20
##   threshold  f1    f2 f0point5 accuracy precision recall specificity
##   <dbl> <dbl> <dbl>   <dbl>   <dbl>   <dbl> <dbl>   <dbl>
## 1    0.324 0.679 0.688   0.671   0.923   0.665 0.695   0.954
## # ... with 12 more variables: absolute_mcc <dbl>, min_per_class_accuracy <dbl>,
## #   mean_per_class_accuracy <dbl>, tns <dbl>, fns <dbl>, fps <dbl>, tps <dbl>,
## #   tnr <dbl>, fnr <dbl>, fpr <dbl>, tpr <dbl>, idx <int>
```

```
performance_tbl %>%
  ggplot(aes(x = threshold)) +
  geom_line(aes(y = precision), color = "blue", size = 1) +
  geom_line(aes(y = recall), color = "red", size = 1) +

  # Insert line where precision and recall are harmonically optimized
  geom_vline(xintercept = h2o.find_threshold_by_max_metric(performance_h2o, "f1")) +
  labs(title = "Precision vs Recall", y = "value") +
  theme_new
```



ROC Plot

```
path <- "/StackedEnsemble_AllModels_AutoML_20210108_015304"

load_model_performance_metrics <- function(path, test_tbl) {

  model_h2o <- h2o.loadModel(path)
  perf_h2o <- h2o.performance(model_h2o, newdata = as.h2o(test_tbl))

  perf_h2o %>%
    h2o.metric() %>%
    as_tibble() %>%
    mutate(auc = h2o.auc(perf_h2o)) %>%
    select(tpr, fpr, auc)

}

#model_metrics_tbl <- fs::dir_info(path = "ml_journal-Automated_Machine_learning_2_model/") %>%
#  select(path) %>%
#  mutate(metrics = map(path, load_model_performance_metrics, test_tbl)) %>%
#  unnest(cols = metrics)

#model_metrics_tbl %>%
#  mutate(
#    # Extract the model names
#    path = str_split(path, pattern = "/", simplify = T)[,2] %>% as_factor(),
```

```

#       auc = auc %>% round(3) %>% as.character() %>% as_factor()
#       ) %>%
#     ggplot(aes(fpr, tpr, color = path, linetype = auc)) +
#     geom_line(size = 1) +
#
#     # just for demonstration purposes
#     geom_abline(color = "red", linetype = "dotted") +
#
#     theme_new +
#     theme(
#       legend.direction = "vertical",
#     ) +
#     labs(
#       title = "ROC Plot",
#       subtitle = "Performance of 3 Top Performing Models"
#     )

```

Precision vs Recall plot

```

# Precision vs Recall

load_model_performance_metrics <- function(path, test_tbl) {

  model_h2o <- h2o.loadModel(path)
  perf_h2o <- h2o.performance(model_h2o, newdata = as.h2o(test_tbl))

  perf_h2o %>%
    h2o.metric() %>%
    as_tibble() %>%
    mutate(auc = h2o.auc(perf_h2o)) %>%
    select(tpr, fpr, auc, precision, recall)

}

#model_metrics_tbl <- fs::dir_info(path = "ml_journal-Automated_Machine_learning_2_model/") %>%
#   select(path) %>%
#   mutate(metrics = map(path, load_model_performance_metrics, test_tbl)) %>%
#   unnest(cols = metrics)
#
#model_metrics_tbl %>%
#   mutate(
#     path = str_split(path, pattern = "/", simplify = T)[,2] %>% as_factor(),
#     auc = auc %>% round(3) %>% as.character() %>% as_factor()
#   ) %>%
#   ggplot(aes(recall, precision, color = path, linetype = auc)) +
#   geom_line(size = 1) +
#   theme_new +
#   theme(
#     legend.direction = "vertical",
#   ) +
#   labs(
#     title = "Precision vs Recall Plot",

```

```
# subtitle = "Performance of 3 Top Performing Models"
# )
```

getting predictions_tbl from previous session

```
predictions <- h2o.predict(stacked_ensemble_h2o, newdata = as.h2o(test_tbl))
```

```
## |
## |
```

```
typeof(predictions)
```

```
## [1] "environment"
```

```
# [1] "environment"
```

```
predictions_tbl <- predictions %>% as_tibble()
```

```
# No    0.938624996 0.06137500
# Yes   0.767573922 0.23242608
# No    0.934670085 0.06532991
# Yes   0.580364130 0.41963587
# Yes   0.738244189 0.26175581
# Yes   0.767517540 0.23248246
# Yes   0.580365130 0.41963487
# No    0.940360179 0.05963982
# Yes   0.616682970 0.38331703
# Yes   0.111431957 0.88856804
```

```
# Gain & Lift
```

```
ranked_predictions_tbl <- predictions_tbl %>%
  bind_cols(test_tbl) %>%
  select(predict:Yes, went_on_backorder) %>%
  arrange(desc(Yes))
```

```
ranked_predictions_tbl
```

```
## # A tibble: 2,857 x 4
##   predict      No   Yes went_on_backorder
##   <fct>      <dbl> <dbl> <fct>
## 1 Yes      0.000523 0.999 Yes
## 2 Yes      0.00168 0.998 Yes
## 3 Yes      0.00280 0.997 Yes
## 4 Yes      0.00334 0.997 Yes
## 5 Yes      0.00348 0.997 Yes
## 6 Yes      0.00388 0.996 Yes
## 7 Yes      0.00397 0.996 Yes
## 8 Yes      0.00435 0.996 Yes
## 9 Yes      0.00437 0.996 Yes
## 10 Yes     0.00449 0.996 Yes
## # ... with 2,847 more rows
```

Gain and Lift calculations

```
ranked_predictions_tbl %>%
  mutate(ntile = ntile(Yes, n = 10)) %>%
  group_by(ntile) %>%
  summarise(
    cases = n(),
    responses = sum(went_on_backorder == "Yes")
  ) %>%
  arrange(desc(ntile))
```

```
## # A tibble: 10 x 3
##   ntile cases responses
##   <int> <int>     <int>
## 1     10     285       200
## 2      9     285        82
## 3      8     285        32
## 4      7     286        12
## 5      6     286         1
## 6      5     286         3
## 7      4     286         3
## 8      3     286         1
## 9      2     286         0
## 10     1     286         0
```

```
calculated_gain_lift_tbl <- ranked_predictions_tbl %>%
  mutate(ntile = ntile(Yes, n = 10)) %>%
  group_by(ntile) %>%
  summarise(
    cases = n(),
    responses = sum(went_on_backorder == "Yes")
  ) %>%
  arrange(desc(ntile)) %>%

  # Add group numbers (opposite of ntile)
  mutate(group = row_number()) %>%
  select(group, cases, responses) %>%

  # Calculations
  mutate(
    cumulative_responses = cumsum(responses),
    pct_responses        = responses / sum(responses),
    gain                 = cumsum(pct_responses),
    cumulative_pct_cases = cumsum(cases) / sum(cases),
    lift                 = gain / cumulative_pct_cases,
    gain_baseline        = cumulative_pct_cases,
    lift_baseline        = gain_baseline / cumulative_pct_cases
  )

calculated_gain_lift_tbl
```

```
## # A tibble: 10 x 10
```



```
##      group cases responses cumulative_resp~ pct_responses  gain cumulative_pct_~
##      <int> <int>      <int>          <int>          <dbl> <dbl>          <dbl>
##  1      1    285        200            200        0.599  0.599          0.0998
##  2      2    285         82            282        0.246  0.844          0.200
##  3      3    285         32            314        0.0958  0.940          0.299
##  4      4    286         12            326        0.0359  0.976          0.399
##  5      5    286          1            327        0.00299  0.979          0.499
##  6      6    286          3            330        0.00898  0.988          0.600
##  7      7    286          3            333        0.00898  0.997          0.700
##  8      8    286          1            334        0.00299  1.00          0.800
##  9      9    286          0            334          0        1.00          0.900
## 10     10    286          0            334          0        1.00          1
## # ... with 3 more variables: lift <dbl>, gain_baseline <dbl>,
## #   lift_baseline <dbl>
```

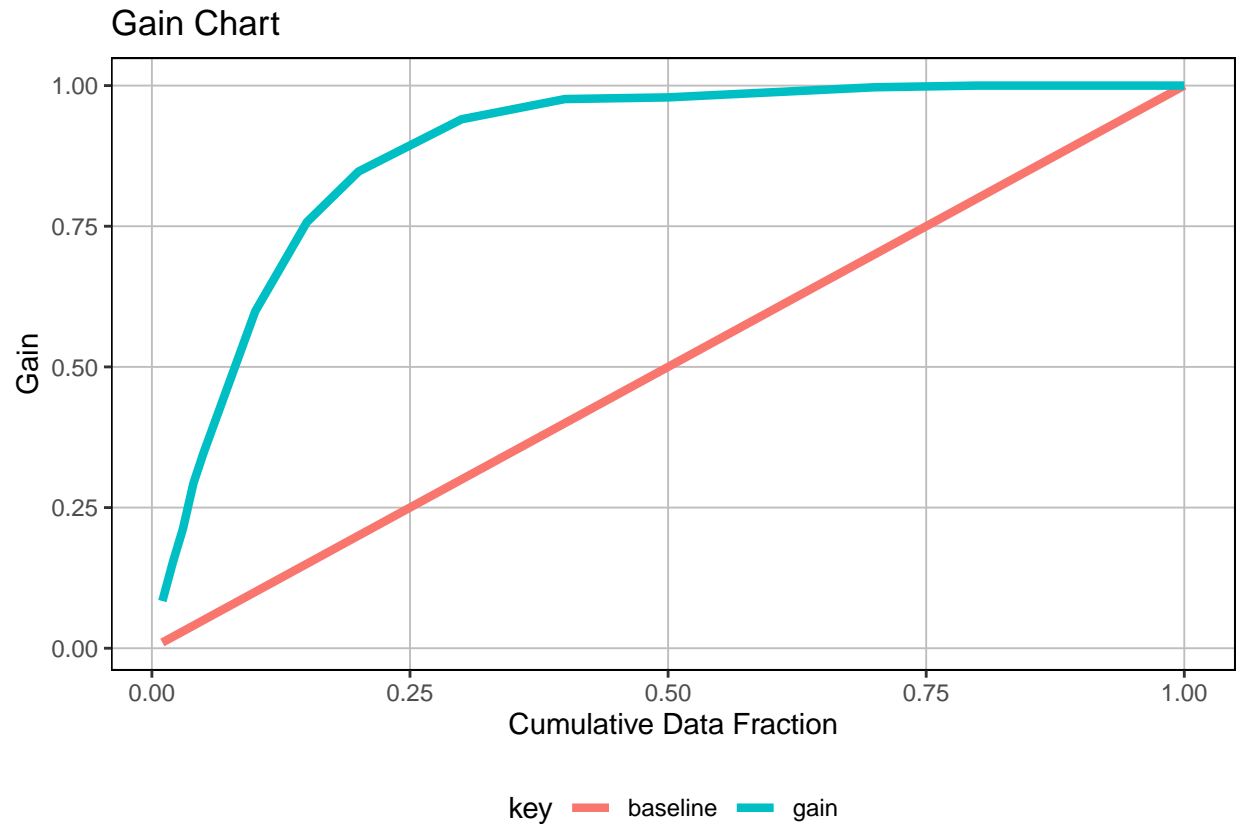
Gain Plot

```
gain_lift_tbl <- performance_h2o %>%
  h2o.gainsLift() %>%
  as.tibble()

## Gain Chart

gain_transformed_tbl <- gain_lift_tbl %>%
  select(group, cumulative_data_fraction, cumulative_capture_rate, cumulative_lift) %>%
  select(-contains("lift")) %>%
  mutate(baseline = cumulative_data_fraction) %>%
  rename(gain      = cumulative_capture_rate) %>%
  # prepare the data for the plotting (for the color and group aesthetics)
  pivot_longer(cols = c(gain, baseline), values_to = "value", names_to = "key")

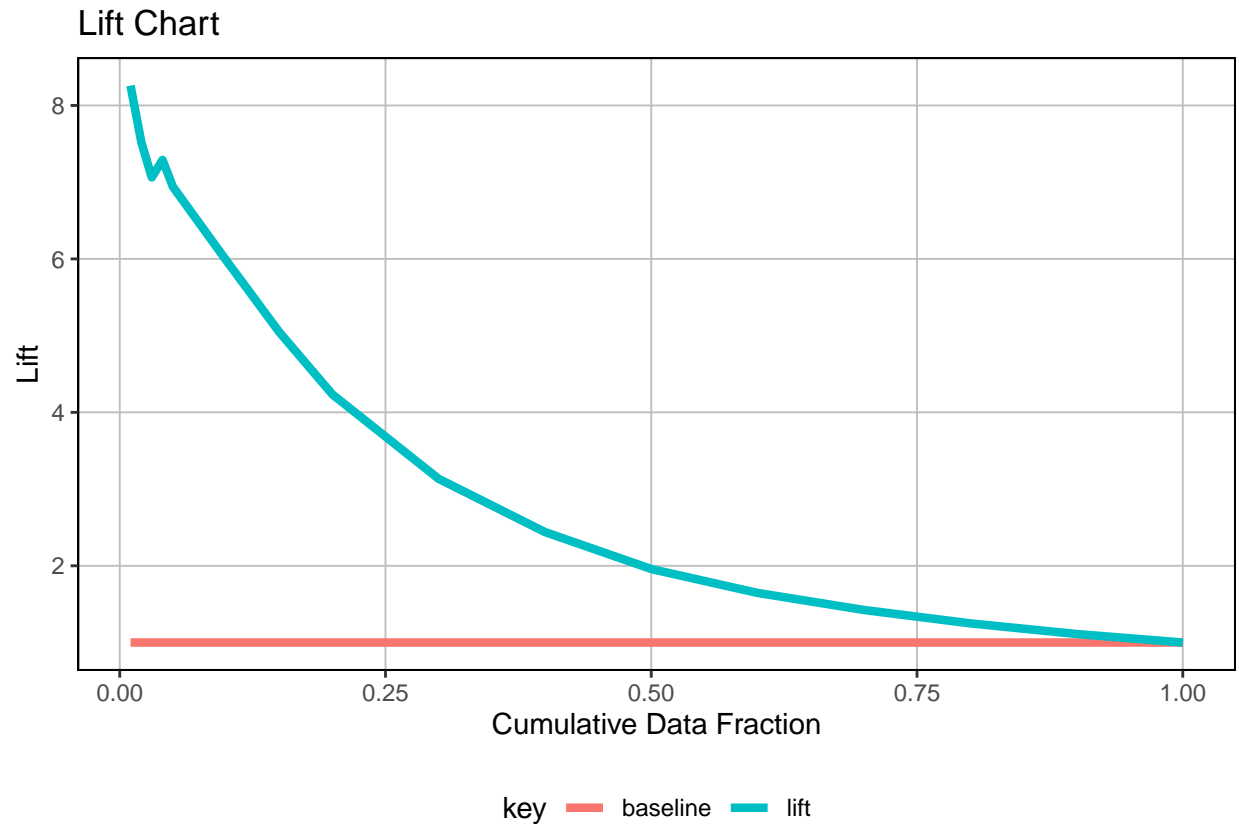
gain_transformed_tbl %>%
  ggplot(aes(x = cumulative_data_fraction, y = value, color = key)) +
  geom_line(size = 1.5) +
  labs(
    title = "Gain Chart",
    x = "Cumulative Data Fraction",
    y = "Gain"
  ) +
  theme_new
```



Lift Plot

```
lift_transformed_tbl <- gain_lift_tbl %>%
  select(group, cumulative_data_fraction, cumulative_capture_rate, cumulative_lift) %>%
  select(-contains("capture")) %>%
  mutate(baseline = 1) %>%
  rename(lift = cumulative_lift) %>%
  pivot_longer(cols = c(lift, baseline), values_to = "value", names_to = "key")

lift_transformed_tbl %>%
  ggplot(aes(x = cumulative_data_fraction, y = value, color = key)) +
  geom_line(size = 1.5) +
  labs(
    title = "Lift Chart",
    x = "Cumulative Data Fraction",
    y = "Lift"
  ) +
  theme_new
```



Cowplot block

Visualization

```
# set values to test the function while building it
h2o_leaderboard <- automl_models_h2o@leaderboard
newdata <- test_tbl
order_by <- "auc"
max_models <- 4
size <- 1

plot_h2o_performance <- function(h2o_leaderboard, newdata, order_by = c("auc", "logloss"),
                                max_models = 3, size = 1.5) {

  # Inputs
  leaderboard_tbl <- h2o_leaderboard %>%
    as_tibble() %>%
    slice(1:max_models)

  newdata_tbl <- newdata %>%
    as_tibble()

  # Selecting the first, if nothing is provided
  order_by <- tolower(order_by[[1]])
```

```

# Convert string stored in a variable to column name (symbol)
order_by_expr <- rlang::sym(order_by)

# Turn off the progress bars ( opposite h2o.show_progress())
h2o.no_progress()

# 1. Model metrics

get_model_performance_metrics <- function(model_id, test_tbl) {

  model_h2o <- h2o.getModel(model_id)
  perf_h2o <- h2o.performance(model_h2o, newdata = as.h2o(test_tbl))

  perf_h2o %>%
    h2o.metric() %>%
    as.tibble() %>%
    select(threshold, tpr, fpr, precision, recall)
}

model_metrics_tbl <- leaderboard_tbl %>%
mutate(metrics = map(model_id, get_model_performance_metrics, newdata_tbl)) %>%
unnest(cols = metrics) %>%
mutate(
  model_id = as_factor(model_id) %>%
    # programmatically reorder factors depending on order_by
    fct_reorder(!! order_by_expr,
                .desc = ifelse(order_by == "auc", TRUE, FALSE)),

  auc = auc %>%
    round(3) %>%
    as.character() %>%
    as_factor() %>%
    fct_reorder(as.numeric(model_id)),

  logloss = logloss %>%
    round(4) %>%
    as.character() %>%
    as_factor() %>%
    fct_reorder(as.numeric(model_id))
)

# 1A. ROC Plot

p1 <- model_metrics_tbl %>%
  ggplot(aes(fpr, tpr, color = model_id, linetype = !! order_by_expr)) +
  geom_line(size = size) +
  theme_new +
  labs(title = "ROC", x = "FPR", y = "TPR") +
  theme(legend.direction = "vertical")

# 1B. Precision vs Recall

p2 <- model_metrics_tbl %>%
  ggplot(aes(recall, precision, color = model_id, linetype = !! order_by_expr)) +

```

```

geom_line(size = size) +
theme_new +
labs(title = "Precision Vs Recall", x = "Recall", y = "Precision") +
theme(legend.position = "none")

# 2. Gain / Lift

get_gain_lift <- function(model_id, test_tbl) {

  model_h2o <- h2o.getModel(model_id)
  perf_h2o <- h2o.performance(model_h2o, newdata = as.h2o(test_tbl))

  perf_h2o %>%
    h2o.gainsLift() %>%
    as.tibble() %>%
    select(group, cumulative_data_fraction, cumulative_capture_rate, cumulative_lift)

}

gain_lift_tbl <- leaderboard_tbl %>%
  mutate(metrics = map(model_id, get_gain_lift, newdata_tbl)) %>%
  unnest(cols = metrics) %>%
  mutate(
    model_id = as_factor(model_id) %>%
      fct_reorder(! order_by_expr,
                  .desc = ifelse(order_by == "auc", TRUE, FALSE)),
    auc = auc %>%
      round(3) %>%
      as.character() %>%
      as_factor() %>%
      fct_reorder(as.numeric(model_id)),
    logloss = logloss %>%
      round(4) %>%
      as.character() %>%
      as_factor() %>%
      fct_reorder(as.numeric(model_id))
  ) %>%
  rename(
    gain = cumulative_capture_rate,
    lift = cumulative_lift
  )

# 2A. Gain Plot

p3 <- gain_lift_tbl %>%
  ggplot(aes(cumulative_data_fraction, gain,
             color = model_id, linetype = ! order_by_expr)) +
  geom_line(size = size,) +
  geom_segment(x = 0, y = 0, xend = 1, yend = 1,
             color = "red", size = size, linetype = "dotted") +
  theme_new +
  expand_limits(x = c(0, 1), y = c(0, 1)) +

```

```

    labs(title = "Gain",
          x = "Cumulative Data Fraction", y = "Gain") +
    theme(legend.position = "none")

# 2B. Lift Plot

p4 <- gain_lift_tbl %>%
  ggplot(aes(cumulative_data_fraction, lift,
             color = model_id, linetype = !! order_by_expr)) +
  geom_line(size = size) +
  geom_segment(x = 0, y = 1, xend = 1, yend = 1,
              color = "red", size = size, linetype = "dotted") +
  theme_new +
  expand_limits(x = c(0, 1), y = c(0, 1)) +
  labs(title = "Lift",
        x = "Cumulative Data Fraction", y = "Lift") +
  theme(legend.position = "none")

# Combine using cowplot

# cowplot::get_legend extracts a legend from a ggplot object
p_legend <- get_legend(p1)
# Remove legend from p1
p1 <- p1 + theme(legend.position = "none")

# cowplot::plt_grid() combines multiple ggplots into a single cowplot object
p <- cowplot::plot_grid(p1, p2, p3, p4, ncol = 2)

# cowplot::ggdraw() sets up a drawing layer
p_title <- ggdraw() +

  # cowplot::draw_label() draws text on a ggdraw layer / ggplot object
  draw_label("H2O Model Metrics", size = 18, fontface = "bold",
            color = "#2C3E50")

p_subtitle <- ggdraw() +
  draw_label(glue("Ordered by {toupper(order_by)}"), size = 10,
            color = "#2C3E50")

# Combine everything
ret <- plot_grid(p_title, p_subtitle, p, p_legend,

               # Adjust the relative spacing, so that the legends always fits
               ncol = 1, rel_heights = c(0.05, 0.05, 1, 0.05 * max_models))

h2o.show_progress()

return(ret)
}

automl_models_h2o@leaderboard %>%

```

```
plot_h2o_performance(newdata = test_tbl, order_by = "logloss",
                     size = 0.5, max_models = 4)
```

H2O Model Metrics

Ordered by LOGLOSS

